Curator: Shravani Bhupal | sbhupal2@illinois.edu

IS 547 - Foundations of Data Curation

Spring 2026

Date - March 30, 2026

Curation Log — SWAN Baseline Dataset (ICPSR 28762)

This is my running record of what I have done, what I found, and any decisions I made along the way. I am keeping this as a plain log rather than a polished document because the point is to show my actual process, including dead ends and things I had to go back and reconsider.

March 3 — Got started. Downloaded the dataset from ICPSR (Study 28762) and set up a folder for the project. First thing I noticed on the study page is that the current version is v5, updated in2019. There is a version history listed on the ICPSR page, but it is not bundled with the files you download, which already struck me as a minor usability issue. I am going to reconstruct the version history manually from the page notes.

March 4 — Downloaded all the files. Here is what was in the deposit:

ICPSR_28762-0001.sav (SPSS data file, ~2.1 MB)

ICPSR_28762-0001.sas7bdat (SAS data file)

ICPSR_28762-0001-Codebook.pdf (247 pages)

ICPSR_28762-0001-User_Guide.pdf (~45 pages)

ICPSR_28762-0001-Setup.sas (SAS syntax)

ICPSR_28762-0001-Setup.sps (SPSS syntax)

Publications_List.pdf (added in v5)

Curator: Shravani Bhupal | sbhupal2@illinois.edu

IS 547 - Foundations of Data Curation

Spring 2026

Date - March 30, 2026

Filter | Cols: « < 1 – 50 > »

	SWANID Respondent ID	VISIT Study visit	INTDAY0 Interview day	AGE0 Age in years	PREGNAN0 Currently pregnant?	PREVBLO0 Blood drawn previously	ALCHL240 Alcohol in last 24 hours	EATDRIN0 Eat/drink in last 12 hours	STRTPER0 Start period in last week
1	10005	00	0	48	1	NA	1	1	2
2	10046	00	0	52	1	NA	1	1	2
3	10056	00	0	51	1	1	2	1	2
4	10092	00	0	45	1	NA	2	1	2
5	10126	00	0	48	1	NA	1	1	2
6	10153	00	0	51	1	1	1	1	2
7	10196	00	0	46	1	NA	1	1	2
8	10245	00	0	47	1	NA	1	1	2
9	10258	00	0	46	1	NA	1	1	2
10	10262	00	0	47	1	NA	1	1	2
11	10386	00	0	44	1	NA	1	1	2
12	10464	00	0	42	1	NA	1	1	2
13	10484	00	0	50	1	1	2	1	2
14	10514	00	0	46	1	NA	1	1	1
15	10522	00	0	46	1	NA	2	1	2
16	10532	00	0	43	1	NA	1	1	2
17	10550	00	0	50	1	NA	1	1	1
18	10604	00	0	45	1	1	1	1	2
19	10623	00	0	43	1	NA	1	1	2
20	10629	00	0	48	1	NA	1	1	1
21	10630	00	0	44	1	NA	1	1	2
22	10694	00	0	44	1	1	2	1	2
23	10747	00	0	46	1	NA	1	1	2
24	10757	00	0	51	1	NA	1	1	2
25	10801	00	0	44	1	NA	1	1	2
26	10879	00	0	44	1	NA	1	1	2
27	10910	00	0	49	1	1	1	1	2
28	10968	00	0	44	1	NA	1	1	2

Showing 1 to 28 of 3,302 entries, 741 total columns

I imported the sav file into RStudio to view the column names with labels to understand the data as there was no direct excel sheet with column labels and variables used.

Curator: Shravani Bhupal | sbhupal2@illinois.edu

IS 547 - Foundations of Data Curation

Spring 2026

Date - March 30, 2026

```
Console   Terminal ×   Background Jobs ×
R 4.4.1 · ~/

R version 4.4.1 (2024-06-14) -- "Race for Your Life"
Copyright (C) 2024 The R Foundation for Statistical Computing
Platform: aarch64-apple-darwin20

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

  Natural language support but running in an English locale

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

[Workspace loaded from ~/.RData]

Loading required package: Matrix
> load("/Users/shravanibhupal/Downloads/ICPSR_28762 6/DS0001/28762-0001-Data.rda")
> load("/Users/shravanibhupal/Downloads/ICPSR_28762 15/DS0001/28762-0001-Data.rda")
> library(haven)
> X28762_0001_Data <- read_sav("Downloads/ICPSR_28762 11/DS0001/28762-0001-Data.sav")
> View(X28762_0001_Data)
> write.csv(X28762_0001_Data, "SWAN_dataset.csv", row.names = FALSE)
>
> getwd()
[1] "/Users/shravanibhupal"
>
> library(haven)
```

No README anywhere in the deposit. I was expecting at least a README based on what good deposits usually look like. Will flag this.

March 5 — Opened the SPSS file. Basic checks:

N = 3,302 — matches what the documentation says

379 variables — matches the codebook

SWANID is unique across all rows, no duplicates

Two kinds of missing values: system missing and -1. The -1 code is not documented consistently across variables, some have it explained in the codebook and some do not

Curator: Shravani Bhupal | sbhupal2@illinois.edu

IS 547 - Foundations of Data Curation

Spring 2026

Date - March 30, 2026

Found a handful of variables where the data has a numeric code that is not listed in the value labels at all — flagging for the gap analysis

The technical side of the deposit is fine. My professor mentioned in his feedback that ICPSR professionally curates deposits like this one, so I was not expecting to find big data errors, and I did not. The interesting curation work is going to be at the documentation level.

March 6-7 — Spent time on the version history. The ICPSR study page lists five versions:

v1 (September 2010): Initial deposit. ICPSR created variable and value labels, checked for undocumented codes.

v2 and v3: Minor updates, limited detail on what changed.

v4 (2018): Missing values were adjusted. One participant's race/ethnicity was corrected —they had been mislabeled.

v5 (May 2019): This is the big update. Variable labels were revised to spell out abbreviations. Variables were reordered to match the PI codebook. Fuller question text was added for many variables. A publications list was added. Online analysis was enabled.

The frustrating thing is that versions v1 through v4 are not available to download, so I cannot actually compare the earlier state of the deposit to what is there now. I noted in my project plant that version history transparency was one of the two main curation concerns, and this confirm sit — a new user downloading the current files has no way to know what changed or when without going to the study page separately.

March 10-13 — Read the full codebook and User Guide. This took a few days because the codebook is 247 pages. Some observations:

The User Guide does a decent job describing the study design, the seven sites, the sampling approach, and the demographic composition of the sample. The racial and ethnic over sampling design is mentioned, but it is described in a way that explains what was done rather than giving users any guidance on how to account for it in their own analysis. Someone new to this dataset could easily miss that issue.

Curator: Shravani Bhupal | sbhupal2@illinois.edu

IS 547 - Foundations of Data Curation

Spring 2026

Date - March 30, 2026

IRB approval is mentioned in the User Guide in general terms. There is no site-specific IRB information — no numbers, no dates, no mention of the coordinating IRB.

The variable labels are noticeably better in v5 than what I can piece together from the change notes about earlier versions. Abbreviations have been expanded, which helps. But question text is still absent for a lot of variables, especially in the employment section.

The employment and occupational stress variables stuck out to me as the weakest part of the codebook. I went in knowing this was going to be a problem area — it is partly why I chose this dataset — and the read confirmed it. That domain is going to be the main focus of the Understand step.

March 16-20 — Variable domain mapping. I went through the full variable list and grouped everything by domain, then gave each domain an initial documentation quality rating based on what the codebook provides.

Domain	Approx. count	Quality
Demographics	~30	Good
Menopausal status	~25	Good
Physical health / biomarkers	~60	Good
Mental health (CES-D, anxiety)	~25	Good
Health behaviors	~20	Good
Medications and medical history	~30	Moderate
Employment and occupational stress	~35	Poor
Social support	~20	Moderate
Site and administrative	~15	Good